

STAT 231: Problem Set 2A

Lorraine Oloo

due by 5 PM on Monday, March 1

In order to most effectively digest the textbook chapter readings – and the new R commands each presents – series A homework assignments are designed to encourage you to read the textbook chapters actively and in line with the textbook’s Prop Tip of page 33:

“Pro Tip: If you want to learn how to use a particular command, we highly recommend running the example code on your own”

A more thorough reading and light practice of the textbook chapter prior to class allows us to dive quicker and deeper into the topics and commands during class. Furthermore, learning a programming language is like learning any other language – practice, practice, practice is the key to fluency. By having two assignments each week, I hope to encourage practice throughout the week. A little coding each day will take you a long way!

Series A assignments are intended to be completed individually. While most of our work in this class will be collaborative, it is important each individual completes the active readings. The problems should be straightforward based on the textbook readings, but if you have any questions, feel free to ask me!

Steps to proceed:

1. In RStudio, go to File > Open Project, navigate to the folder with the course-content repo, select the course-content project (course-content.Rproj), and click "Open"
2. Pull the course-content repo (e.g. using the blue-ish down arrow in the Git tab in upper right window)
3. Copy ps2A.Rmd from the course repo to your repo (see page 6 of the GitHub Classroom Guide for Stat231 if needed)
4. Close the course-content repo project in RStudio
5. Open YOUR repo project in RStudio
6. In the ps2A.Rmd file in YOUR repo, replace "YOUR NAME HERE" with your name
7. Add in your responses, committing and pushing to YOUR repo in appropriate places along the way
8. Run "Knit PDF"
9. Upload the pdf to Gradescope. Don’t forget to select which of your pages are associated with each problem. *You will not get credit for work on unassigned pages (e.g., if you only selected the first page but your solution spans two pages, you would lose points for any part on the second page that the grader can’t see).*

1. NYC Flights

a.

In Section 4.3.1, the `flights` and `carrier` tables within the `nycflights13` package are joined together. Recreate the `flightsJoined` dataset from page 80. Hint: make sure you've loaded the `nycflights13` package before referring to the data tables (see code on page 79).

```
library(nycflights13)
#glimpse(flights)

flights_joined <- flights %>%
  inner_join(airlines, by = c("carrier" = "carrier"))
glimpse(flights_joined)

## Rows: 336,776
## Columns: 20
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 849,~
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851,~
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -1~
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA",~
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD",~
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6~
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-01 0~
## $ name      <chr> "United Air Lines Inc.", "United Air Lines Inc.", "Amer~
```

b.

Now, create a new dataset `flightsJoined2` that:

- creates a new variable, `distance_km`, which is distance in kilometers (note that 1 mile is about 1.6 kilometers)
- keeps only the variables: `name`, `flight`, `arr_delay`, and `distance_km`
- keeps only observations where distance is less than 500 kilometers

Hint: see examples in Section 4.1 for subsetting datasets and creating new variables.

```
select(flights_joined2 <- flights_joined %>%
  mutate(distance_km = distance*1.6), name, flight, arr_delay, distance_km)
```

```
## # A tibble: 336,776 x 4
##   name                flight arr_delay distance_km
##   <chr>                <int>      <dbl>      <dbl>
## 1 United Air Lines Inc.   1545         11      2240
## 2 United Air Lines Inc.   1714         20     2266.
## 3 American Airlines Inc.  1141         33     1742.
## 4 JetBlue Airways        725        -18     2522.
## 5 Delta Air Lines Inc.    461        -25     1219.
## 6 United Air Lines Inc.  1696         12     1150.
## 7 JetBlue Airways        507         19     1704
## 8 ExpressJet Airlines Inc. 5708        -14       366.
## 9 JetBlue Airways         79         -8     1510.
## 10 American Airlines Inc.  301          8     1173.
## # ... with 336,766 more rows
```

c.

Lastly, using the functions introduced in Section 4.1.4, compute the number of flights (call this *N*), the average arrival delay (call this *avg_arr_delay*), and the average distance in kilometers (call this *avg_dist_km*) among these flights with distances less than 500 km (i.e. working off of *flightsJoined2*) *grouping by the carrier name*. Sort the results in descending order based on *avg_arr_delay*.

Getting NAs for *avg_arr_delay*? That happens when some observations are missing that data. Before grouping and summarizing, add a line to exclude observations with missing arrival delay information using *filter(is.na(arr_delay)==FALSE)*.

```
flights_joined3<- flights_joined2 %>%
  select(name, flight, arr_delay, distance_km) %>%
  filter(distance_km<500 & is.na(arr_delay)==FALSE)

fj<- flights_joined3 %>%
  group_by(name) %>%
  summarize(
    N = sum(flight),
    avg_arr_delay = mean(arr_delay),
    avg_dist_km = mean(distance_km)
  )

arrange(fj, desc(avg_arr_delay))
```

```
## # A tibble: 11 x 4
##   name                N avg_arr_delay avg_dist_km
##   <chr>                <int>      <dbl>      <dbl>
## 1 Mesa Airlines Inc.   1071507      18.0       360.
## 2 ExpressJet Airlines Inc. 70431260     15.6       373.
## 3 Envoy Air           10207234     11.0       351.
## 4 JetBlue Airways     12555455      8.66       385.
## 5 Endeavor Air Inc.    22179923      6.82       339.
```

| | | | | |
|----|---------------------------|----------|--------|------|
| ## | 6 Southwest Airlines Co. | 289201 | 4.92 | 272. |
| ## | 7 United Air Lines Inc. | 3164913 | 4.09 | 320. |
| ## | 8 SkyWest Airlines Inc. | 4967 | 3 | 366. |
| ## | 9 US Airways Inc. | 19331098 | 2.22 | 308. |
| ## | 10 American Airlines Inc. | 1999580 | 1.88 | 299. |
| ## | 11 Delta Air Lines Inc. | 1679867 | -0.643 | 325. |

2. Baby names

a.

Working with the `babynames` data table in the `babynames` package, create a dataset `babynames2` that only includes years 2000 to 2017.

```
data = babynames::babynames

babynames2<- babynames::babynames %>%

  filter(year %in% 2000:2017)

babynames2
```

```
## # A tibble: 591,925 x 5
##   year sex   name      n    prop
##   <dbl> <chr> <chr>   <int> <dbl>
## 1  2000 F     Emily   25953 0.0130
## 2  2000 F    Hannah  23080 0.0116
## 3  2000 F   Madison 19967 0.0100
## 4  2000 F   Ashley 17997 0.00902
## 5  2000 F    Sarah 17697 0.00887
## 6  2000 F   Alexis 17629 0.00884
## 7  2000 F  Samantha 17266 0.00866
## 8  2000 F   Jessica 15709 0.00787
## 9  2000 F Elizabeth 15094 0.00757
## 10 2000 F    Taylor 15078 0.00756
## # ... with 591,915 more rows
```

b.

Following the code presented in Section 5.2.4, create a dataset called `BabyNarrow` that summarizes the total number of people with each name (born between 2000 and 2017), grouped by sex. (Hint: follow the second code chunk on page 102, but don't filter on any particular names.) Look at the dataset. Why have we called this dataset “narrow”?

ANSWER: in the `sex` column, one gender is listed first, the next gender is listed next instead of two different columns for the M and F.

```
BabyNarrow<- babynames2 %>%
group_by(sex,name) %>%
summarise(total = sum(n))
```

```
## 'summarise()' has grouped output by 'sex'. You can override using the '.groups' argument.
```

```
BabyNarrow
```

```
## # A tibble: 73,332 x 3
## # Groups:   sex [2]
##   sex   name      total
```

```
##      <chr> <chr>      <int>
##  1 F      Aabha        35
##  2 F      Aabriella    32
##  3 F      Aada         5
##  4 F      Aaden        5
##  5 F      Aadhira      77
##  6 F      Aadhvika     9
##  7 F      Aadhya     1478
##  8 F      Aadi         16
##  9 F      Aadilynn     5
## 10 F      Aadison      11
## # ... with 73,322 more rows
```

c.

Now, following the code chunk presented on page 103*, put the data into a wide format (call the new dataset `BabyWide`), and only keep observations where both M and F are greater than 10,000. Compute the `ratio` (as `pmin(M/F, F/M)`) and identify the top three names with the largest ratio. (Note: these names could be different from the ones found on page 103 since we limited the dataset to years 2000-2017 and names with greater than 10,000 individuals.)

- Note: you can use the `pivot_wider()` function instead of the `spread()` function if using the 2nd edition of the textbook (e.g., see Section 6.2.2 and 6.2.3 in the 2nd edition). I find `pivot_wider()` and `pivot_longer()` to be more intuitive than `spread()` and `gather()`.

ANSWER:

Blair
Elisha
Kerry

```
# this will bring up "Pivoting Introduction" vignette in your Help tab
#vignette("pivot")
BabyWide <- babynames2 %>%

group_by(sex, name) %>%
summarize(total = sum(n)) %>%
spread(key = sex, value = total, fill = 0)
```

'summarise()' has grouped output by 'sex'. You can override using the '.groups' argument.

BabyWide

```
## # A tibble: 67,063 x 3
##   name      F      M
##   <chr>    <dbl> <dbl>
##  1 Aaban      0    107
##  2 Aabha     35      0
##  3 Aabid      0     10
##  4 Aabir      0      5
##  5 Aabriella  32      0
```

```
## 6 Aada      5      0
## 7 Aadam     0     202
## 8 Aadan     0     130
## 9 Aadarsh   0     199
## 10 Aaden    5    4653
## # ... with 67,053 more rows
```

How to combine the 2 r chunks??

```
BabyWide %>%
  filter(M > 10000, F > 10000) %>%
  mutate(ratio = pmin(M / F, F / M) ) %>%
  arrange(desc(ratio)) %>%
  head(3)
```

```
## # A tibble: 3 x 4
##   name      F      M ratio
##   <chr>  <dbl> <dbl> <dbl>
## 1 Justice 10947 11267 0.972
## 2 Skyler  17120 22154 0.773
## 3 Quinn   25022 19080 0.763
```

d.

Lastly, use the `gather()` function (or the `pivot_longer()` function) to put the dataset back into narrow form. Call this dataset `BabyNarrow2`. Hint: see Section 5.2.3. Why are the number of observations in `BabyNarrow2` different from that in `BabyNarrow`?

ANSWER: The filter function used in the making of `BabyWide` where both `F` and `M` were > 10000

```
BabyNarrow2 <- BabyWide %>%
  gather(key = sex, value = total, F, M)

BabyNarrow2
```

```
## # A tibble: 134,126 x 3
##   name      sex  total
##   <chr>  <chr> <dbl>
## 1 Aaban    F        0
## 2 Aabha    F       35
## 3 Aabid    F        0
## 4 Aabir    F        0
## 5 Aabriella F       32
## 6 Aada     F        5
## 7 Aadam    F        0
## 8 Aadan    F        0
## 9 Aadarsh  F        0
## 10 Aaden   F        5
## # ... with 134,116 more rows
```

BabyNarrow

```
## # A tibble: 73,332 x 3
## # Groups:   sex [2]
##   sex  name      total
##   <chr> <chr>    <int>
## 1 F    Aabha      35
## 2 F    Aabriella  32
## 3 F    Aada       5
## 4 F    Aaden       5
## 5 F    Aadhira    77
## 6 F    Aadhvika    9
## 7 F    Aadhya   1478
## 8 F    Aadi       16
## 9 F    Aadilynn   5
## 10 F   Aadison   11
## # ... with 73,322 more rows
```